

HomeActivity: Recognizing home activities using sensor data

Stephen Lee, Patrick Pagus and Dong Chen

December 16, 2015

1 Abstract:

Home activities recognizing allows many potential smart applications including healthcare and energy efficiency. In this project, we present the state of the art models for recognizing activities, and show how they performance on the Kasteren-a real world dataset. Comparative analysis metrics including precision, recall, accuracy, F1 mesure and MCC are used in the experiments. Our evaluation results could be used for other pattern recognizing area rather than home activity recognizing.

2 Introduction

Automatically recognizing human activities such as cooking, sleeping and bathing, in a home setting allows many applications in areas such as intelligent environments and healthcare. Recent developments in sensing technology has led to wireless sensor networks which provide a non-intrusive, privacy friendly and easy to install solution to in-home monitoring. Sensors used are contact switches to measure open-close states of doors and cupboards; pressure mats to measure sitting on a couch or lying in bed; mercury contacts for movement of objects such as drawers; passive infrared sensors to detect motion in a specific area and float sensors to measure the toilet being flushed.

Activity recognition is the problem of determining an individual's or group's activities given some sensor data describing their actions. More formally, activity recognition finds a function that maps a sequence of sensor readings (feature vectors), to a sequence of activity labels that best describes the sensor data. We focus on activity recognition within home based wireless sensor networks.

In this case, sensors may include contact, motion, temperature, and humidity sensors located on doors, appliances or monitoring certain rooms. This problem has many applications including healthcare and energy efficiency, where other sensors, such as cameras and wearables, intrude upon privacy or are too expensive or inconvenient. In healthcare, home activity recognition can be used to assess the cognitive and physical capabilities of an elderly person.

Home energy efficiency can be improved given a schedule of the occupant’s activities by scheduling background loads, such as heating, cooling, and various appliances. This domain has several challenges. First, there is a dearth of training data because labeling one’s activities is a tedious time consuming process, this limits the number of effective techniques. Second, patterns in training data are specific to the home and the individual. Therefore, the performance of a machine learning technique may vary widely over different data sets. Finally, dominant classes may confound machine learning techniques.

To further our knowledge of machine learning, we will conduct a comprehensive empirical analysis of proposed methods in this domain, across a collection of datasets. These methods include Hidden Markov Models(HMM), Conditional Random Fields(CRF), and Support Vector Machines(SVM). We will apply these methods to at least four datasets that vary in duration, home size, and sensor counts, which are not necessarily proportional to home size. By applying these techniques to real-world datasets and explaining the differences and similarities in their performances, we will gain a deeper understanding of machine learning.

For our group project, we tackle the problem of activity recognition using available sensor information installed in homes. These sensors could be infrared motion detector sensors (to understand when a person entered a room or opened a fridge), to RFID sensors, to microphone based labeling of activities to record the ground truth. These datasets are usually available in the form of timestamp and the corresponding activity or sensor label. So, in this project we would like to predict a label (i.e. the activity) given a sequence of observations.

3 Related Works

The probabilistic models discussed in our project represent the state of the art models used in activity recognition. Tapia et al. used the naive Bayes model in combination with the raw feature representation on two real world datasets recorded using a wireless sensor network [1]. HMMs were used in work by Patterson et al. and were applied to data obtained from a wearable RFID reader in a house where many objects are equipped with RFID tags [2]. In work by van Kasteren et al. the performance of HMMs and CRFs in activity recognition was compared on a realworld dataset recorded using a wireless sensor network [3]. Duong et al. compared the performance of HSMMs and HMMs in activity recognition using a laboratory setup in which four cameras captured the location of a person in a kitchen setup [4]. One type of model that we have not included in our comparison are hierarchical models. They have been successfully applied to activity recognition from video data [5], in an office environment using cameras, microphones and keyboard input [6] and on data obtained from a wearable sensing system [7]. The related works show part of the models in their work, respectively. It is hard to compare across different works. Instead, we implemented a set of models and ran the experiments on the same real word dataset.

The evaluation results could be used as the baseline for future research.

4 Research Problem

We use more models, compared with other related works.

5 Recognizing Methods

We would like to explore the following techniques:

1. HMM - Hidden Markov Modes (HMM) are a ubiquitous tool for modeling time series data and a generative probabilistic model consisting of a hidden variable and an observable variable at each time step. HMM can be used as black-box density models on sequences. They have the advantage over Markov models in that they can represent long-range dependencies between observations, mediated via the latent variables.

In our project the hidden variable is the activity performed, and the observable variable is the vector of sensor readings. There are two dependency assumptions: The hidden variable at time t , namely y_t , depends only on the previous hidden variable y_{t-1} . The observable variable at time t , namely x_t , depends only on the hidden variable y_t at that time slice. Then, we can specify an HMM using three probability distributions: the distribution over initial states $p(y_1)$; the transition distribution $p(y_t | y_{t-1})$ representing the probability of going from one state to the next; and the observation distribution $p(x_t | y_t)$ indicating the probability that the state y_t would generate observation x_t . Learning the parameters of these distributions corresponds to maximizing the joint probability $p(x, y)$ of the paired observation and label sequences in the training data. We can factorize the joint distribution in terms of the three distributions described above as follows:

$$P(x, y) = \prod_{t=1}^T p(y_t | y_{t-1}) p(x_t | y_t), p(y_1) = p(y_1 | y_0) \quad (1)$$

The parameters that maximize this joint probability are found by frequency counting. Because in our case we are dealing with discrete data, we can simply count the number of occurrences of transitions, observations and states [15]

2. CRF - A Conditional Random Field (CRF) is a discriminative probabilistic model that can come in many different forms. The form that most closely resembles the HMM is known as a linear-chain CRF and is the model we use in this project.

$$p(y \mid x) = \frac{1}{Z(x)} \exp \left[\sum_{k=1}^K \lambda_k f_k(y_t, y_{t-1}, x_t) \right] \quad (2)$$

$$Z(x) = \sum_y \exp \left[\sum_{k=1}^K \lambda_k f_k(y_t, y_{t-1}, x_t) \right] \quad (3)$$

3. SVM - The implementation used in this experiment comes from SVM
4. SSVM -
5. Naive Bayes

$$p(y_{1:T}, x_{1:T}) = \prod_{t=1}^T p(x_t \mid y_t) p(y_t) \quad (4)$$

$$p(x_t \mid y_t = i) = \prod_{n=1}^N p(x_t^n \mid y_t = i) \quad (5)$$

$$p(x_t^n \mid y_t = i) = \mu_{ni}^{x_t^n} (1 - \mu_{ni})^{1-x_t^n} \quad (6)$$

The datasets available to us are:

1. Kasteren dataset - has over a month long sensor information from 3 homes.
2. Tulum dataset - more than six month period from a single home

6 Experimental Evaluation

6.1 Datasets

The Kasteren dataset is recording a 26-year-old man. He lives alone in a three-room apartment where 14 state-change sensors were installed. Locations of sensors include doors, cup-boards, refrigerator and a toilet flush sensor. Sensors were left unattended, collecting data for 28 days in the apartment. This resulted in 2120 sensor events and 245 activity instances.

6.2 Experimental Setup

6.3 Feather Representation

Raw

Changepoint

Last-fired

6.4 Comparative Analysis Metrics

$$Precision = \frac{TP}{TP + FP} \quad (7)$$

$$Recall = \frac{TP}{TP + FN} \quad (8)$$

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \quad (9)$$

$$F - Measure = \frac{2 \cdot Precision \cdot Recall}{Precision + Recall} \quad (10)$$

The Matthews correlation coefficient is used in machine learning as a measure of the quality of binary (two-class) classifications, introduced by biochemist Brian W. Matthews in 1975. It takes into account true and false positives and negatives and is generally regarded as a balanced measure which can be used even if the classes are of very different sizes. The MCC is in essence a correlation coefficient between the observed and predicted binary classifications; it returns a value between -1 and +1. A coefficient of +1 represents a perfect prediction, 0 no better than random prediction and -1 indicates total disagreement between prediction and observation.

$$MCC = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP + FP) \cdot (TP + FN) \cdot (TN + FP) \cdot (TN + FN)}} \quad (11)$$

6.5 Experiments

7 Discussion and Future Work

8 Reference

References

- [1] E. M. Tapia, S. S. Intille, and K. Larson. *Activity recognition in the home using simple and ubiquitous sensors*. In Pervasive Computing, Second International Conference, PERVASIVE 2004, pp. 158175, Vienna, Austria (April, 2004).




Figure 1: House A

- [2] D. J. Patterson, D. Fox, H. A. Kautz, and M. Philipose. Fine-grained activity recognition by aggregating abstract object usage. In ISWC, pp. 4451. IEEE Computer Society, (2005). ISBN 0-7695-2419-2. URL <http://doi.ieeecomputersociety.org/10.1109/ISWC.2005.22>.
- [3] T. van Kasteren, A. Noulas, G. Englebienne, and B. Krose. Accurate activity recognition in a home setting. In UbiComp 08: Proceedings of the 10th international conference on Ubiquitous computing, pp. 19, New York, NY, USA, (2008). ACM. ISBN 978-1-60558-136-1. doi: <http://doi.acm.org/10.1145/1409635.1409637>.
- [4] T. Duong, D. Phung, H. Bui, and S. Venkatesh, Efficient duration and hierarchical modeling for human activity recognition, *Artif. Intell.* 173(7-8), 830856, (2009). ISSN 0004-3702. doi: <http://dx.doi.org/10.1016/j.artint.2008.12.005>.
- [5] S. Luhr, H. H. Bui, S. Venkatesh, and G. A. West, Recognition of human activity through hierarchical stochastic learning, *percom.* 00, 416, (2003). doi: <http://doi.ieeecomputersociety.org/10.1109/PERCOM.2003.1192766>.




Figure 2: House B

- [6] N. Oliver, A. Garg, and E. Horvitz, Layered representations for learning and inferring office activity from multiple sensory channels, *Comput. Vis. Image Underst.* 96(2), 163180, (2004). ISSN 1077-3142. doi: <http://dx.doi.org/10.1016/j.cviu.2004.02.004>.
- [7] A. Subramanya, A. Raj, J. Bilmes, and D. Fox. Hierarchical models for activity recognition. In *IEEE Multimedia Signal Processing (MMSP) Conference*, Victoria, CA (October, 2006).




Figure 3: House C